

CS 2112—Spring 2014

Final Project

CritterWorld

Version of February 19, 2014

The final four programming assignments for this course make up a single project that you will work on as part of a two-person team. In this project, you will simulate a simple world of animals (“critters”!) that wander around, eat food, reproduce, and evolve. This world will include a graphical visualization that enables a user to take control of individual animals. Different species can also interact in this simulated world. Because of evolution, there will be multiple species eventually even if initially there was only one.

Critters move around on a regular, hexagonal grid. Critters need energy to survive, because everything they do requires energy. They gain a little energy from the sun each time step as long as they are not moving. Critters may also attack each other; when critters run out of energy, they die and become food for other critters. If critters accumulate enough energy, they can reproduce.

The genome of a critter is actually an event-driven program that determines what the critter does on each time step. When critters reproduce, the genome is copied from the parent or parents to the new critter, with possibly some mutations applied. This means that critter programs may change over time and perhaps evolve to make them more effective.

The critter simulation will be implemented as a networked Java service with a graphical front end. This design permits multiple users to view and interact with the same virtual world.

In summary, this project involves developing the following components:

- a simple parser and interpreter for the critter language
- a graphical user interface
- a distributed implementation of the system

0 Changes to the spec

The following changes have been made since the initial release.

- No changes yet.

1 The world

The world that critters live in is a large array of hexagonal tiles called *hexes*. A given critter is at any moment located on one of these hexes and facing in one of the six possible directions, as shown in Figure 1. The world advances in discrete time steps. On each time step, a critter may perform one of several possible actions, or simply update its internal state and wait for the next time step while absorbing solar energy.




Figure 1: Hexes and directions




Figure 2: Rows and columns of hexes

Hexes are either rock hexes or hexes that contain other things. Rocks are inert obstacles that critters cannot move onto or over. A non-rock hex may be empty, or it may contain a critter, some food, or a critter and some food. A hex cannot contain more than one critter.

A certain amount of food is created on a hex when a critter dies there. Other critters may then eat the food to gain energy.

Hex coordinates. Each hex is identified by a coordinate (c, r) where c is the column and r is the row on which the hex is located. Both c and r are nonnegative integers. Figure 2 shows the row and column coordinates of various hexes. The world has a fixed, roughly rectangular shape, in which each column contains the same number of hexes. The lower-left (southwest corner) hex is at coordinate $(0,0)$. Moving in one of the six possible directions changes the row and column coordinate of the critter. The corresponding adjustments to row and column coordinates are shown in Figure 1.

Some coordinates lie outside the world. A coordinate that lies outside the world acts in all ways as though it is a rock hex. Critters cannot fall off the edge of the world, and they see rock when they look off the edge. Figure 2 shows a very small world with $\text{MAX_COLUMN} = 5$ and $\text{MAX_ROW} = 7$. A coordinate such that $c \not\leq \text{MAX_COLUMN}$ lies off the east edge of the world. A coordinate such as $(1,0)$, where $c \not\leq 2r$, lies off the south edge of the world. And a coordinate where $c + (2 \cdot \text{MAX_ROW} - \text{MAX_COLUMN}) \not\leq 2r$ lies off the north edge of the world. The world will be roughly square if MAX_ROW is about 1.37 times as large as MAX_COLUMN .

You might notice that there is a third potential “coordinate” we could specify, corresponding to lines of hexes slanting upward toward the right. We can define the “slice” s of a given hex as $(r - c)$, increasing as we move upward along a column or “northwest” along a row. Hexes in the lower right corner of the map have a negative slice. Using slices, we can compactly represent the

length of the shortest path between two hexes:

$$\max \{|\Delta c|, |\Delta r|, |\Delta s|\} = \max \{|\Delta c|, |\Delta r|, |\Delta r - \Delta c|\}$$

Time The simulation proceeds in time steps. During each time step, each critter is allowed to take one turn. These turns are taken sequentially, so each critter sees the changes to the state of the world caused by all critters that have already taken a turn during the current time step. The order in which critters take turns is fixed. Newly created critters are added to the end of the ordering. However, being early in the ordering is not an advantage because nothing happens to the state of the world between time steps.

2 Critter actions

The allowed actions are the following:

- **Wait.** The critter waits until the next turn without doing anything except absorbing solar energy.
- **Move forward or backward.** A critter uses some energy to move forward to the hex in front of it or backward to the hex behind it. If it attempts to move and there is a critter or rock in the destination hex, the move fails (but still takes energy)
- **Turn.** It may rotate 60 degrees right or left. This takes little energy.
- **Eat.** A critter may eat as much food that is sitting on its current hex as it can. It gains energy from any food that it eats. If the amount of food that is on the hex is more than it can absorb, the remaining food is left on the hex.
- **Serve.** A critter may convert some of its own energy into food dropped onto its hex.
- **Attack.** It may attack a critter directly in front of it. The attack removes an amount of energy from the attacked critter that is determined by the size and offensive ability of the attacker and the defensive ability of the victim.
- **Tag.** The critter may tag the critter in front of it—for example, to mark that critter as an enemy or friend.
- **Grow.** A critter may use energy to increase its size by one unit.
- **Bud.** A critter may use a large amount of its energy to produce a new, smaller critter with the same genome (possibly with some random mutations).
- **Mate.** A critter may attempt to mate with another critter in front of it. For this to be successful, the critter in front must also be facing toward it and attempting to mate in the same time step. If mating is successful, both critters use energy to create new size 1 critter whose genome is the result of merging the genomes of its parents. Unsuccessful mating uses little energy.

3 Critter state

Critter state comprises several attributes, along with the program that drives the critter. The critter has a current location in the world and a current direction, represented as an integer in 0..5, as shown in Figure 1. It also has a current size and energy. The critter also has some fixed attributes: its offensive and defensive ability, the size of its memory, and the rules governing its behavior.

Two parts of critter state control how it appears to other critters: its *posture*, which it can change, and its *tag*, which other critters can change. The posture is an integer between 0 and 999, which the critter may change arbitrarily. The tag is an integer between 0 and 999 and is initially 0. It may be changed by any other critter using its tag action, but a critter is not able to change its own tag.

Each critter has a derived attribute, its *complexity*. This is a weighted sum of the number of its rules, the size of its memory, and its offensive and defensive abilities. The energy of certain actions the critter performs depends on its complexity. The formula for complexity is found in Section 12.

4 Critter memory

Each critter has a memory called *mem*, which is an array of fixed length containing integers. The first few entries in this array have a meaning that is the same for all critter species:

- *mem*[0]: the length of the critter's memory (immutable, always at least 8)
- *mem*[1]: defensive ability (immutable, ≥ 1)
- *mem*[2]: offensive ability (immutable, ≥ 1)
- *mem*[3]: size (variable, but cannot be assigned directly, ≥ 1)
- *mem*[4]: energy (variable, but cannot be assigned directly, ≥ 1)
- *mem*[5]: pass number, explanation below (variable, but cannot be assigned directly, ≥ 1).
- *mem*[6]: tag (variable, but cannot be assigned directly. Always between 0 and 99.)
- *mem*[7]: posture (assignable only to values between 0 and 99).

There are three kinds of memory entries: immutable entries that never change, variable entries that reflect the current state of the critter but that the critter's rules cannot assign to, and general-purpose mutable entries that can be both read from and assigned to by critter rules.

The size of the critter's memory must be at least 8 to accommodate these entries. If the size is larger, the remaining entries, indexed starting at 8, are general-purpose entries.

5 Rule language

The grammar for the rules is given as a context-free grammar in so-called EBNF (Extended Backus-Naur Form¹) in Figure 3. In EBNF grammars, the right-hand side may be a regular expression. EBNF does not add any real expressive power to context-free grammars, but makes them

¹Despite its name, EBNF is apparently not due to either Backus or Naur!

```

program → rule*
      rule → condition --> command ;
command → update* update-or-action
update-or-action → update | action
      update → mem [ expr ] := expr
      action → wait | forward | backward | left | right
              | eat | attack | grow | bud | mate
              | tag [ expr ] | serve [ expr ]
condition → conjunction ( or conjunction )*
conjunction → relation ( and relation )*
      relation → expr rel expr | { condition }
      rel → < | <= | = | >= | > | !=
      expr → term ( addop term )*
      term → factor ( mulop factor )*
      factor → ⟨number⟩ | mem [ expr ] | ( expr ) | sensor
sensor → nearby [ expr ] | ahead [ expr ] | random [ expr ]
addop → + | -
mulop → * | / | mod

```

Figure 3: Grammar for critter rules

easier to express concisely. Terminal symbols are shown in typewriter font. Terminal symbols whose lexical representation is not fixed are shown using angle brackets: for example, ⟨*number*⟩. Non-terminals are shown in italics, like this: *program*.

5.1 Syntactic sugar

Convenient abbreviations may be used in place of expressions of the form *mem*[*n*] for certain literal constant integers *n*. This *syntactic sugar* (so-called because it “sweetens the syntax”) may be used both when programs are read from a file, and when programs are displayed to the user. Regardless of how the expression is written in the input file or displayed to the user, the underlying abstract syntax tree node is always of the form *mem*[*n*]. The full list of abbreviations is as follows:

| Abbreviation | AST representation |
|--------------|--------------------|
| MEMSIZE | mem[0] |
| DEFENSE | mem[1] |
| OFFENSE | mem[2] |
| SIZE | mem[3] |
| ENERGY | mem[4] |
| PASS | mem[5] |
| TAG | mem[6] |
| POSTURE | mem[7] |

Therefore, a rule like `mem[3]>1000 --> mem[11] := mem[11] - mem[4];` can be written and displayed completely equivalently as `SIZE>1000 --> mem[11] := mem[11] - ENERGY;`

5.2 Executing rules

When it is a critter's turn, it finds the first rule in its list of rules whose condition is true. It then performs all of the updates on the right-hand-side of the rule, along with the action, if any. If the command for the rule contains no action, the process repeats: starting again from the very first rule, it finds the earliest rule whose condition holds, and performs its command. This process is performed up to 999 times, after which the critter automatically performs a `wait` action. If on any pass through the rules, no rule's condition is true, the critter's turn immediately ends and it performs a `wait` action on that turn.

The special memory location `mem[5]` (PASS) reports which pass through the rules is being done. It has the value 1 on the first pass through the rules, 2 on the second pass (if any), then 3, and so on up to a maximum of 999. It starts over again at 1 on the critter's next turn.

It may be possible to accelerate running the rules in various ways. For example, later passes might only check rules whose condition could possibly have become true. Or, if the selected rule has no effect on critter state, the critter will never select any other rule and there is no reason to run further passes. However, these sorts of optimizations are not required.

6 Sensing

A critter can sense its immediate surroundings using *sensor* expressions as described in the grammar.

- The expression `nearby[dir]` reports the contents of the hex in direction *dir*, where $0 \leq dir \leq 5$. Here the direction is relative to the critter's current orientation, so 0 is always immediately in front, 1 is 60 degrees to the right, and so on. (If *d* is out of bounds, its remainder when divided by 6 is used.) The contents are reported as a number *n*, as follows:
 - 0: the hex is completely empty
 - $n > 0$: the hex contains a critter with appearance *n* (see Section 7).
 - $n < -1$: the hex contains no critter, and total food value is $-(n + 1)$.

- $n = -1$: the hex contains a rock.
- The expression `ahead[dist]` reports the contents of hex that is directly ahead of the creature at distance *dist*, using the same scheme as `nearby`. If *dist* is not positive, `ahead` evaluates as if it were `ahead[-dist - 1]`, except that critter that is on the hex is ignored. Thus, `ahead[0]` reports on the appearance of the current critter, whereas `ahead[-1]` reports on the hex under the critter, ignoring the critter itself, and `ahead[-2]` reports on the hex ahead of the critter, while ignoring any other critter that is located there.
- The `random` expression generates a random integer from 0 up to one less than the value of the given expression. Thus, `random[2]` gives either 0 or 1 randomly. For $n < 2$, `random[n]` always is zero.

7 Critter appearance

Memory entry `mem[6]` (sugar: TAG) contains the critter's tag, which is initially zero and is set to some other value only when some other critter tags it. The action `tag[expr]`, where *expr* evaluates to some value *v*, causes the critter in front of it to acquire the tag *v*. The action has no effect if *v* is not a legal tag value.

The entry `mem[7]` (sugar: POSTURE) contains the critter's posture, which defines part of how it looks to other critters. The posture can be used as a way to signal to other nearby critters what the current critter is up to, or as a way of signaling species identity. Initially zero, the posture it is set by simply assigning to its memory location.

When a critter is seen by another critter (or by itself, using `ahead[0]`), its appearance is reported as a positive integer, equal to $tag * 100,000 + size * 1,000 + posture * 10 + direction$. Note that both the critter's tag and posture are less than 100. A newly created critter has size 1, posture 0, and tag 0, so its appearance is $1,000 + direction$. Here, *direction* is, as usual, one of $\{0, \dots, 5\}$, and is reported relative to the observing critter's own direction. For example, a direction of 3 means that the observed critter is facing in the direction opposite from that of the observing critter. Critter programs can use the operators `mod` and `/` to extract the different components of the appearance.

8 Energy and size

A critter has an initial size of 1 but can increase its size using the action `grow`. Size affects energy expenditure but also makes the critter more effective at some actions. Size also determines the maximum energy of the critter. For each point of size, the critter can hold `ENERGY_PER_SIZE = 500` points of energy. Any updates that would increase energy beyond this point cause excess energy to be discarded.

If energy ever goes to (or below) zero, the critter dies. Its death adds to the food on its hex a number of food points equal to `FOOD_PER_SIZE (= 200)` points per point of critter size.

9 Attacking and defending

When one critter attacks another, some damage is done to the defending critter (the victim). This subtracts energy from the victim. If the victim's energy goes to zero (or lower), the victim dies and is turned into an amount of food proportional to its size.

When one critter attacks another, the damage done depends on the sizes of the two critters. If critter 1 attacks critter 2, and S_1 and S_2 are the sizes of the corresponding critters, O_1 is the offensive ability of critter 1, and D_2 is the defensive ability of critter 2, the energy removed from critter 2 is:

$$\text{BASE_DAMAGE} \cdot S_1 \cdot P(\text{DAMAGE_INC} \cdot (S_1 \cdot O_1 - S_2 \cdot D_2))$$

where $\text{BASE_DAMAGE} = 100$, $\text{DAMAGE_INC} = 0.2$, and $P(x)$ is the *logistic function*:

$$P(x) = \frac{1}{1 + e^{-x}}$$

This formula means that critters do damage proportional to their size, but that they do only half their maximum damage if they are evenly matched against the defending critter. Damage falls off quickly to zero when attacking a critter with a higher effective defense.

10 Mutation

When a critter's genome is copied to a new critter, the copy may be perfect. But with probability $p = 1/4$ there will be at least one mutation. If there is one mutation, there is then a 1/4 chance (that is, 1/16 overall) of a second mutation, and so on for possible additional mutations.

A mutation is either a change to an attribute or a change to the rule set, with each equally probable. The attributes that may change are the size of the memory and the offensive and defensive abilities. A change to an attribute is an increment or decrement, chosen with equal probability, to one of these three attributes, chosen with equal probability. However, changes to attributes never reduce them below their minimal legal value (8 for memory size, 1 for offense and defense).

A mutation to the rule set is performed by randomly picking a node in the abstract syntax tree describing the entire set of rules. All nodes are chosen with equal probability. Given that a node has been selected, one of the following changes is made, with equal probability among each of the possible alternatives:

1. The node is removed. If its parent node needs a replacement node, one of its children of the right kind is used. The child to be used is randomly selected. Thus, rule nodes are simply removed, but binary operation nodes would be replaced with either their left or right child.
2. The order of two children of the node is switched. For example, this allows swapping the positions of two rules, or changing $a - b$ to $b - a$.
3. The node and its children are replaced with a copy of another randomly selected node of the right kind, found somewhere in the rule set. The entire AST subtree rooted at the selected node is copied.

4. The node is replaced with a randomly chosen node of the same kind (for example, replacing attack with eat, or + with *), but its children remain the same. Literal integer constants are adjusted up or down by the value of `java.lang.Integer.MAX_VALUE/r.nextInt()`, where legal, and where `r` is a `java.util.Random` object.
5. A newly created node is inserted as the parent of the mutated node. The previous parent becomes the parent of the inserted node, and the mutated node becomes a child of the inserted node. If the inserted node has more than one child, the children that are not the original node are copies of randomly chosen nodes of the right kind from the whole rule set.
6. For nodes with a variable number of children, an additional copy of one of the children, randomly selected, is appended to the end of the list of children. This applies to the root node, where a new rule can be added, and also to command nodes, where the sequence of updates can be extended with another update.

11 Budding and mating

When a new critter is created by budding, it appears directly behind the critter doing the budding. When two critters mate, it appears directly behind one of the two critters, chosen at random.

When a new critter is created by budding, its rules are copied from its parent, modulo possible mutation. Its attributes are also copied from the parent, with the exception of energy, size, posture, and tag. Energy is set to a constant `INITIAL_ENERGY = 250`, size is always set to 1, posture is always set to 0, and the tag is always set to 0. All memory locations at or above index 8 are set to zero in the newly created critter.

When two critters mate, however, they exchange genetic material to form the new critter. Attributes 0–2 are chosen from between the two critters randomly. Attributes 3–5 are chosen as for budding. The new rule sequence is chosen by picking the corresponding rule in sequence from either the ‘mother’ or the ‘father’, at random. Thus, the new critter inherits, in general, some rules from each parent. If the mother or father have different-sized rule sets, the new rule set either has the size of the mother or the father, randomly chosen. Thus, if the mother and father have identical genomes, and there are no mutations, the child will have the same genome too.

12 Energy

Different actions take different amounts of energy, even waiting for a turn. The energy cost of different actions is as follows:

- wait : this action *increases* the critter’s energy by its own size times `SOLAR_FLUX (= 1)`.
- tag, turn, eat : energy equal to the critter’s size.
- forward and backward: energy equal to the critter’s size times `MOVE_COST` where `MOVE_COST = 3`.

- *serve* : the amount of energy served onto the hex, plus the creature's size. A critter can use the action *serve* to send its own energy down to zero, killing it, but not below. A critter that kills itself in this way deposits additional food onto the hex in the usual way.
- *attack* : energy equal to the critter's size times $\text{ATTACK_COST} = 5$.
- *grow* : energy equal to $\text{size} \cdot \text{complexity} \cdot \text{GROW_COST}$ where $\text{GROW_COST} = 1$.
- *bud* : $\text{BUD_COST} \cdot \text{complexity}$ energy, where $\text{BUD_COST} = 9$.
- *mate* : $\text{MATE_COST} \cdot \text{complexity}$ energy, where $\text{MATE_COST} = 5$.

Several of these energy costs depend on the critter complexity. If r is the number of rules in the critter program and *offense* and *defense* are the critter's offensive and defensive abilities, the critter complexity is equal to:

$$r \cdot \text{RULE_COST} + (\text{offense} + \text{defense}) \cdot \text{ABILITY_COST}$$

Most actions take the same energy whether they are successful or unsuccessful. One exception is the *mate* action, which only costs as much as turn if it is unsuccessful.

13 Handling out-of-bounds arguments

One important principle is that syntactically legal critter programs always evaluate successfully. Even when an argument to a sensor or action might seem to be out of bounds, the expression or action will complete. Mutation to critter programs can never cause the simulation as a whole to fail.

Handling ostensibly out-of-bounds arguments is handled in the following way for the various language constructs:

- $\text{mem}[\text{expr}]$: A read from a memory location *expr* where *expr* is not a valid memory index always returns 0. An update to an illegal memory location has no effect. An update to a memory location whose value is constrained (in particular, $\text{mem}[7]$) also has no effect if the value is out of bounds for that location.
- $\text{tag}[\text{expr}]$: If *expr* evaluates to a tag value that is illegal, it has no effect.
- $+$ and $-$: these operate exactly like Java $+$ and $-$.
- $/$ and mod : If the divisor is zero, the result of the expression is also zero.
- $\text{nearby}[\text{expr}]$: the remainder of *expr* when divided by 6 is used as the direction.
- $\text{random}[\text{expr}]$: always zero when $\text{expr} < 2$.

14 Example critter program

The following critter program should be able to survive, find food, and reproduce.

```
POSTURE != 17 --> POSTURE := 17; // we are species 17!
{nearby[3] = 0 or nearby[3] < 0-1} and ENERGY > 2500 --> bud;
{ENERGY > SIZE * 400 and SIZE < 7} --> grow;
ahead[0] < 0-1 and ENERGY < 500 * SIZE --> eat;
// next line attacks only other species
(ahead[1] mod 1000) != 17 and ahead[1] > 0 --> attack;
ahead[1] < 0-5 --> forward;
ahead[2] < 0-10 and ahead[1] = 0 --> forward;
ahead[3] < 0-15 and ahead[1] = 0 --> forward;
ahead[4] < 0-20 and ahead[1] = 0 --> forward;
nearby[0] > 0 and {nearby[3] = 0 or nearby[3] < 0-1} --> backward;
random 3 = 1 --> left;
1 = 1 --> wait; // mostly soak up the rays
```

15 Constants

Numbers used in this document are mostly parameters that have symbolic names. We may fiddle with the values of these parameters to make the simulation more interesting, so you should always use the symbolic names rather than hard-coding them into your program. We are providing [a file containing the current values of these constants](#); your program should parse the file at run time to set them to the correct values.

The current values of the simulation constants are shown in Figure 4. Most of the constants are integers, but a few are real numbers, as indicated by the presence of a decimal point.

16 Challenges

This project has several very different kinds of subsystems. One of the major challenges will be to keep the different parts of the system separate, so that, for example, your simulation code is entirely separate from your graphics code. This particular separation will be crucial for making an distributed version of the program, since the simulation and the displays will be done on a completely different machines. Therefore the simulation code needs to avoid knowing about or naming the graphics code. Similarly, you will want to separate different parts of the simulation into different modules. Your programming tasks will be simpler if the interpretation of critter rules is kept separate from the mechanics of the world and even of the critter itself.

Thoughtful design up front along with your partner will save you a tremendous amount of time later on. Meet early with your partner and decide on how you will structure your project and agree on interfaces and specifications that connect the different parts of the code.

| Name | Value | Description |
|--------------------|-------|--|
| BASE_DAMAGE | 100 | The multiplier for all damage done by attacking |
| DAMAGE_INC | 0.2 | Controls how quickly increased offensive or defensive ability affects damage |
| ENERGY_PER_SIZE | 500 | How much energy a critter can have per point of size |
| FOOD_PER_SIZE | 200 | How much food is created per point of size when a critter dies |
| MAX_SMELL_DISTANCE | 10 | Maximum distance at which food can be sensed |
| ROCK_VALUE | -1 | The value reported when a rock is sensed |
| MAX_COLUMN | 49 | Maximum column index in the world map |
| MAX_ROW | 67 | Maximum row index in the world map |
| MAX_RULES_PER_TURN | 999 | The maximum number of rules that can be run per critter turn |
| SOLAR_FLUX | 1 | Energy gained from sun by doing nothing |
| MOVE_COST | 3 | Energy cost of moving (per unit size) |
| ATTACK_COST | 5 | Energy cost of attacking (per unit size) |
| GROW_COST | 1 | Energy cost of growing (per size and complexity) |
| BUD_COST | 9 | Energy cost of budding (per unit complexity) |
| MATE_COST | 5 | Energy cost of successful mating (per unit complexity) |
| RULE_COST | 2 | Complexity cost of having a rule |
| ABILITY_COST | 25 | Complexity cost of having an ability point |
| INITIAL_ENERGY | 250 | Energy of a newly birthed critter |
| MIN_MEMORY | 8 | Minimum number of memory entries in a critter |
| MAX_SMELL_DIST | 10 | Maximum distance at which food can be smelled |

Figure 4: Constants

17 Extensions

You may add extensions to the critter simulation, but you need not. Chances are you'll have your hands full just implementing the project as is, so be judicious about adding new features.

Possible extensions might be additions to the critter language (abbreviations? function definitions?), to the critter model (better sensory capabilities?), changes to the world (volcanos? water? plants that grow food? climate gradients?), better user control over the world view (zooming and panning? high-level critter commands?). Feel free to be creative. If your extensions might interfere with our testing, for example by making our critter programs invalid, be sure to support a command-line flag `-compatible` that turns off your extra features. We recommend being backward compatible to our specification in any case.

18 Tournament

We'll have a tournament at the end of the class when you can bring in some critter programs to compete in various events such as survival, food gathering, and a maze race. We encourage participation in the tournament, and there will be free food, but the fun is optional. We'll post more information about the tournament as it approaches.